

# Decremental Connectivity in Trees

---

Philip Bille

# Outline

---

- Decremental Connectivity in Trees Problem
- Decremental Connectivity in Paths
  - First Tradeoffs
  - Two-Level Solution
- Decremental Connectivity in Trees
  - Two-Level Solution

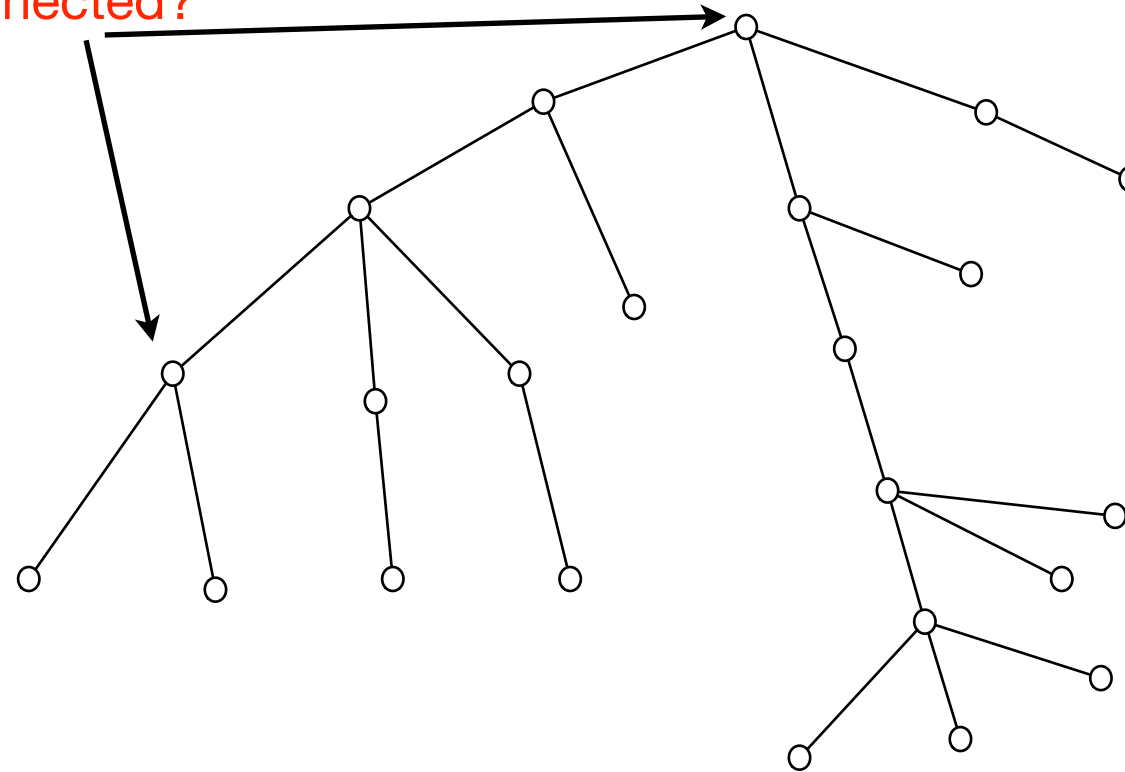
# Decremental Connectivity Problem

# Decremental Connectivity Problem

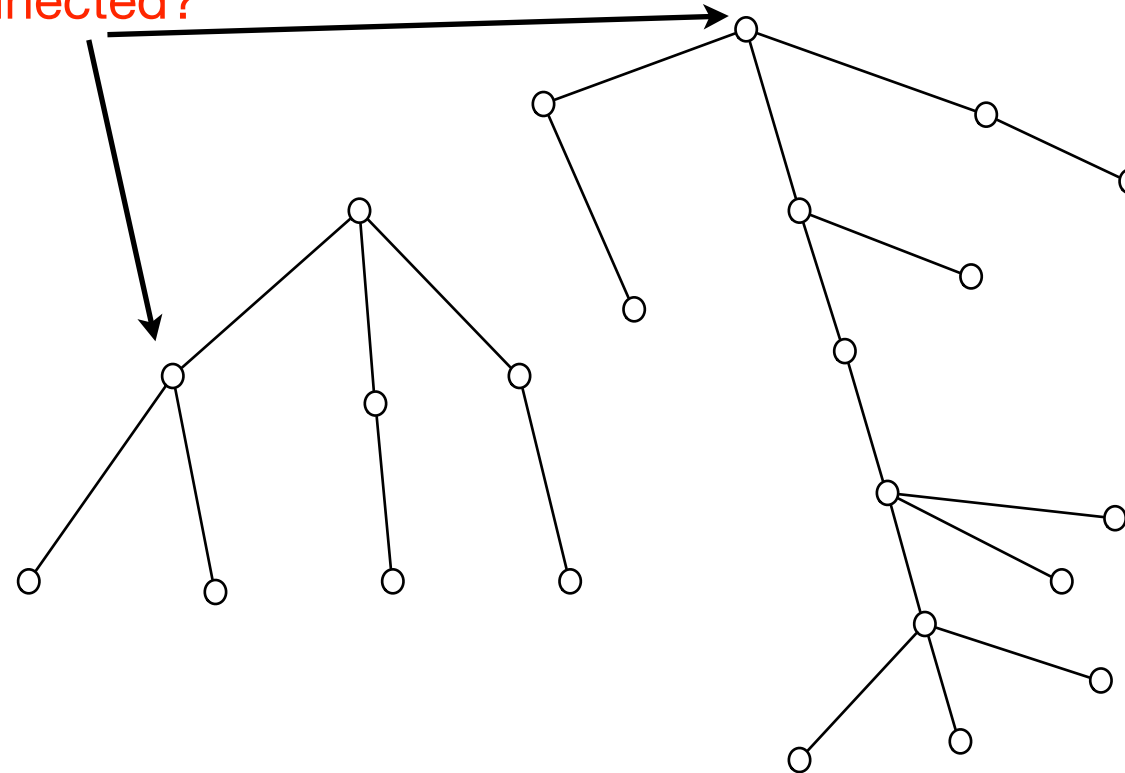
---

- The *decremental connectivity problem*: Starting with a tree  $T$  with  $n$  nodes support the following operations:
- $\text{connected}(v,u)$ : return true if  $v$  and  $w$  are connected.
- $\text{delete}(e)$ : delete the edge  $e$ .

connected?



connected?



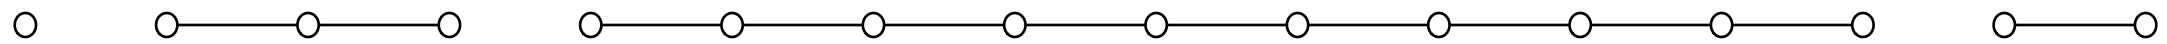
# Applications

---

- Special case of dynamic graph algorithm (inserting and deleting edges/nodes) while supporting some query/queries.
- Nice illustration of techniques for trees and word-level parallelism.
- Nice illustration of algorithmic theory useful in practice.

# Overview

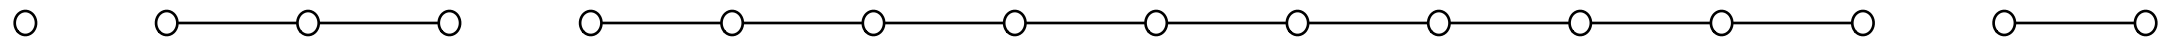
---



- First consider the simple case of *paths*. Later generalize to trees.
- Goal:
  - $O(1)$  for connected.
  - $O(n)$  total for executing all  $n-1$  delete operations.

# Overview

---



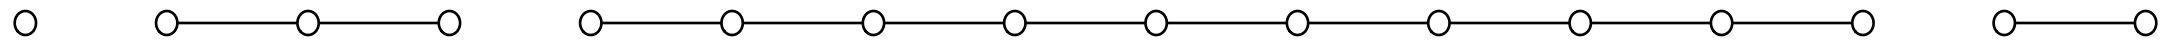
- Solution in 3 steps:
  - First tradeoffs: queries vs. updates
  - Balanced relabeling
  - Clustering with word-level parallelism



# First Tradeoffs

# First Tradeoffs

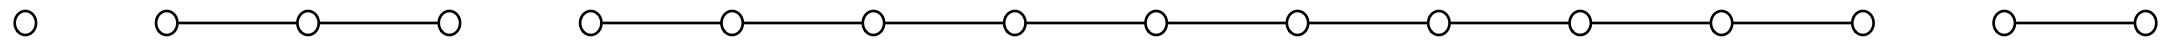
---



- What tradeoffs can we get for connected queries vs. deletions?

# Fast Deletions

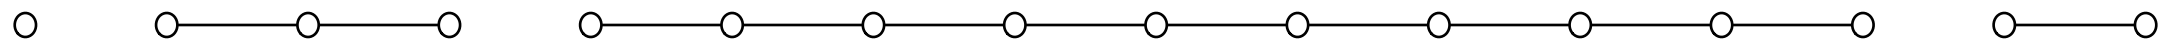
---



- How fast deletions can we get if we ignore the time for connected?

# Fast Deletions

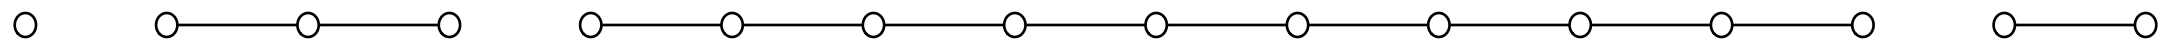
---



- Solution:
  - `delete(e)`: remove `e` from path.
  - `connected(v,u)`: traverse component from `v` to see if `u` is in component containing `v`.
- $O(1)$  for deletion  $\Rightarrow O(n)$  for  $n-1$  deletions
- $O(n)$  for `connected`

# Fast Connected Queries

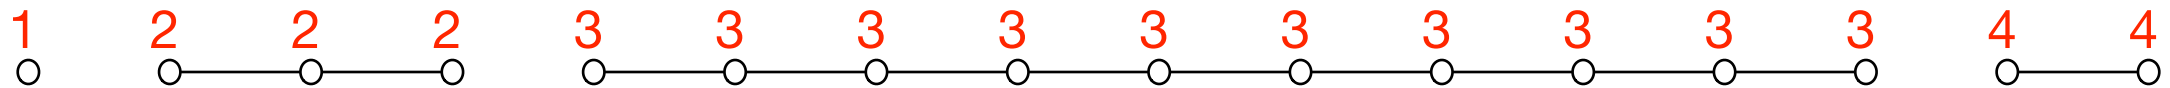
---



- How fast connected queries can we get if we ignore the time for deletions?

# Fast Connected Queries

---



- Solution: Maintain a component ID for each node.
- $v$  and  $u$  are connected iff  $ID(v) = ID(u)$ .
  - $connected(v,u)$ : return true iff  $ID(v) = ID(u)$
  - $delete(e)$ : relabel node IDs of left endpoint of  $e$  (or right endpoint of  $e$ ).
- $O(1)$  for connected
- $O(n)$  for single delete  $\Rightarrow O(n^2)$  for  $n-1$  deletions

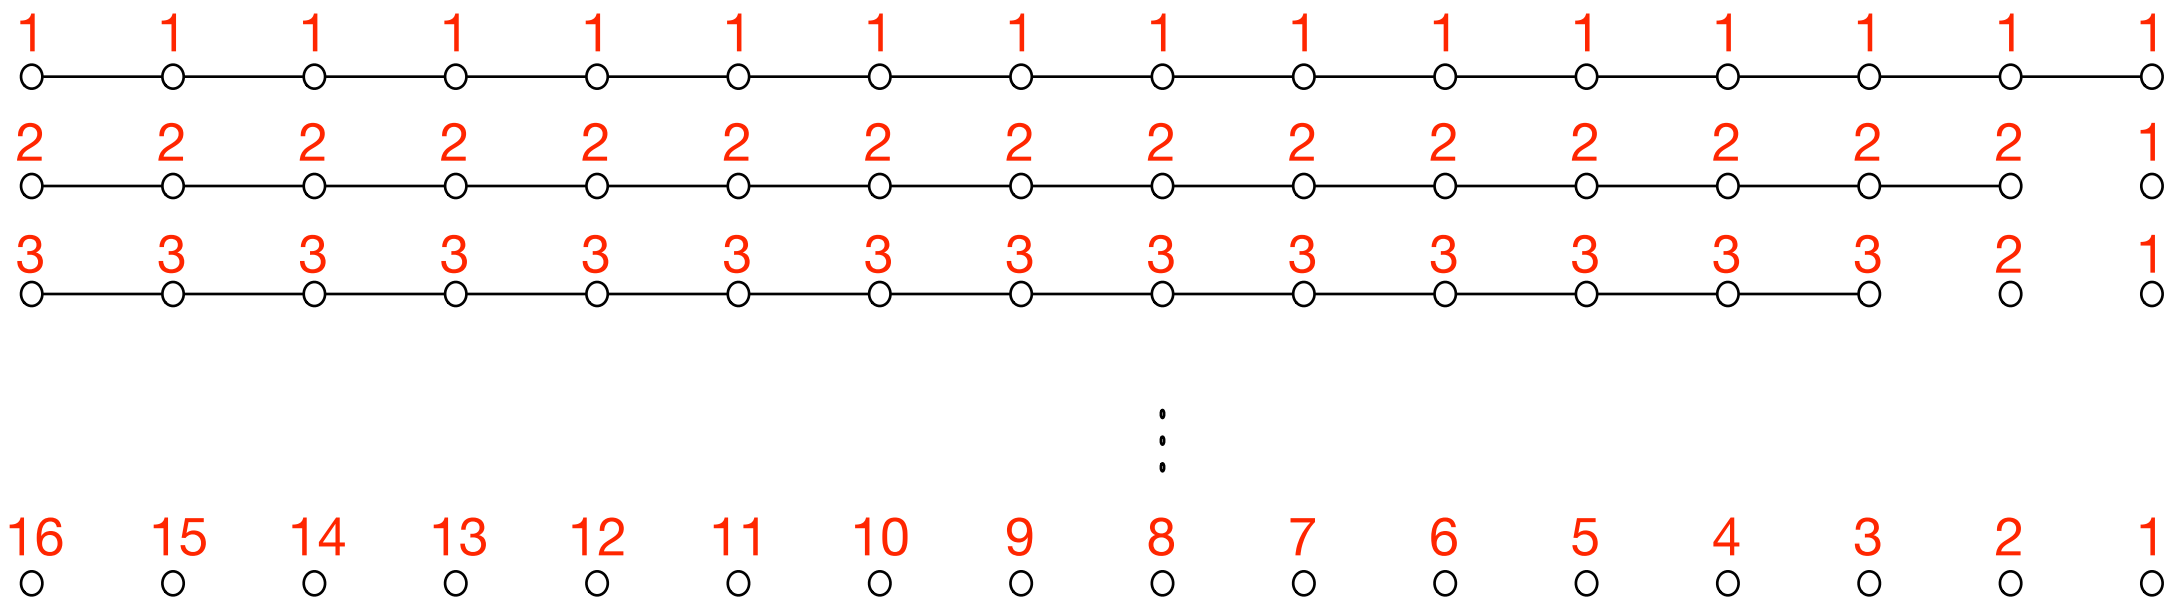
# Better Tradeoffs?

---

- We have two extreme tradeoffs:
  - connected  $O(n)$  and  $n-1$  deletes  $O(n)$ .
  - connected  $O(1)$  and  $n-1$  deletes  $O(n^2)$ .
- How can we get better tradeoffs?
- Consider the bad case for relabeling algorithm.

# A Bad Delete Sequence

---

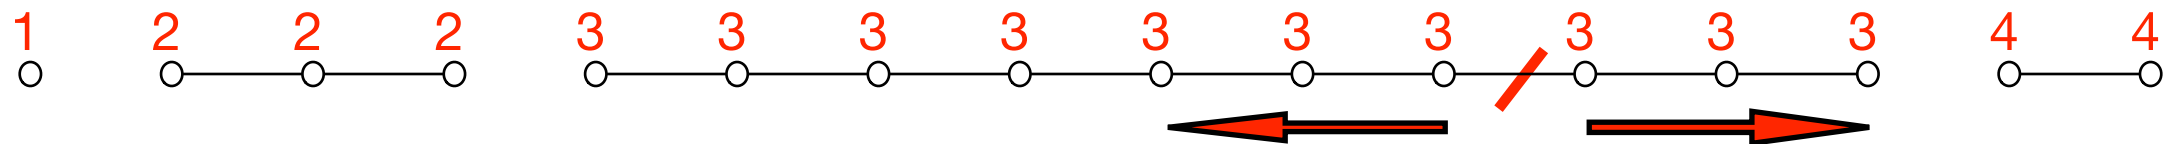


- Main problem:  $n-1$  delete operations take
  - $O((n-1) + (n-2) + \dots + 1) = O(n^2)$
- How can we do better?
- Idea: After a delete update ID for the *smaller component*.



# Balanced Relabeling

---



- New delete(e):
  - Traverse components for endpoints of e *in parallel*.
  - Stop when we find the *smallest* component.
  - Update ID for smallest component.
- ID(v) is only updated when v is in the smaller component.
- => After first update to ID(v) the size of v's component is  $\leq n/2$ , next  $\leq n/4$ , next  $\leq n/8, \dots$
- => At most  $\log n$  updates for ID(v).
- Total time for  $n-1$  deletions =  $O(\text{Total number of ID updates}) = O(n \log n)$ .

# Summary

---

- Theorem: We can solve decremental connectivity in paths in
  - $O(1)$  time for connected.
  - $O(n \log n)$  time for  $n-1$  deletes.
- How can we *shave* off the log-factor?

Shaving a Log-Factor



# LOG FACTORS

*"A goal is a dream with a deadline"*

— Napoleon Hill

# Overview

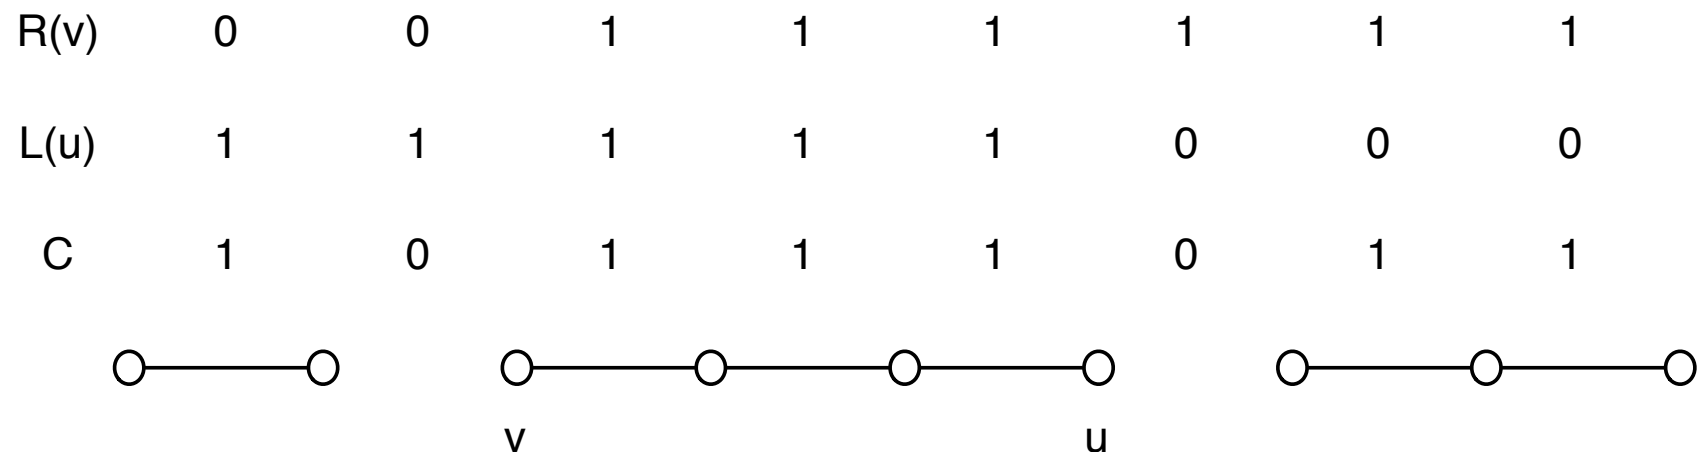
---

- Goal:
  - $O(1)$  time connected
  - $O(n)$  time for  $n-1$  deletes.
- Solution by two-level data structure:
  - Divide path into  $n/w$  subpath of  $w$  nodes.
  - Level 1: Balanced relabeling data structure over path of  $n/w$  nodes.
  - Level 2: New data structure for paths of length  $\leq w$  supporting delete and connected in  $O(1)$  time.

# Data Structure for Short Paths

# Data Structure for Short Paths

---

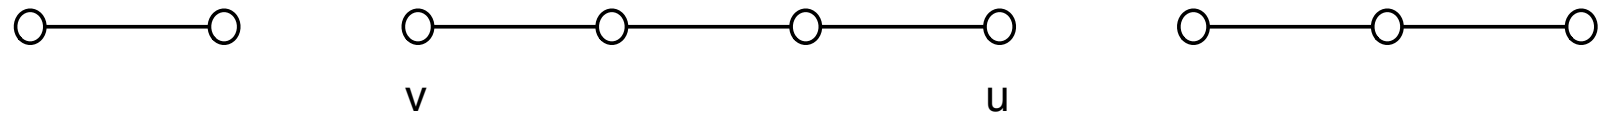


- Paths of length  $\leq w$ .
- Data structure consists of bitmasks:
- C:  $C[i] = 1$  iff edge  $i$  exists.
- For each node  $v$ :
  - R(v):  $R(v)[i] = 1$  iff  $i$  is to the right of  $v$ .
  - L(v):  $L(v)[i] = 1$  iff  $i$  is to the left of  $v$ .
- $O(w)$  bitmasks of length  $\leq w \Rightarrow O(w)$  space.

# Delete

---

R(v)	0	0	1	1	1	1	1	1
L(u)	1	1	1	1	1	0	0	0
C	1	0	1	1	1	0	1	1



- How can we implement delete in  $O(1)$  time?
- delete(edge  $i$ ): set  $C[i] = 0$ .


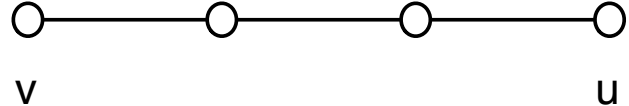



# Connected

---

R(v)	0	0	1	1	1	1	1	1
L(u)	1	1	1	1	1	0	0	0
C	1	0	1	1	1	0	1	1

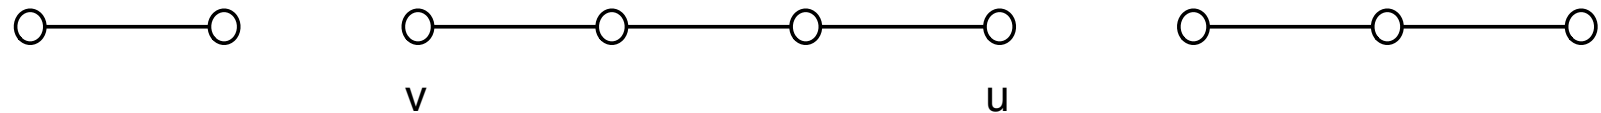
		
---	---	---

- How can we implement connected in  $O(1)$  time?
- $\text{connect}(v,u) = \text{Return true iff } R(v) \ \& \ L(u) \ \& \ \neg C = 0$

# Summary

---

R(v)	0	0	1	1	1	1	1	1
L(u)	1	1	1	1	1	0	0	0
C	1	0	1	1	1	0	1	1

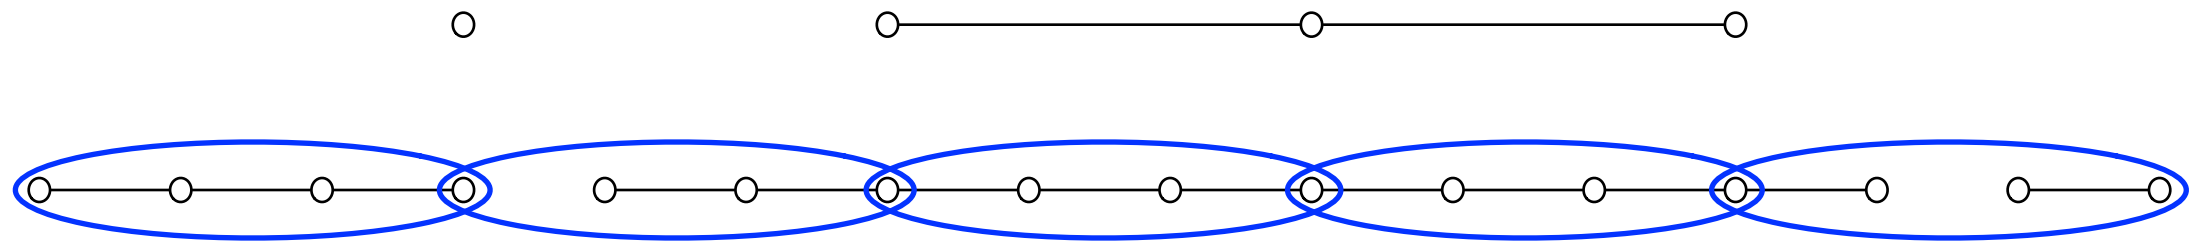


- Theorem: We can solve decremental connectivity in paths of *length*  $\leq w$  in
  - $O(1)$  time for connected.
  - $O(1)$  time for delete.

# Two-Level Data Structure

# A Two-Level Solution

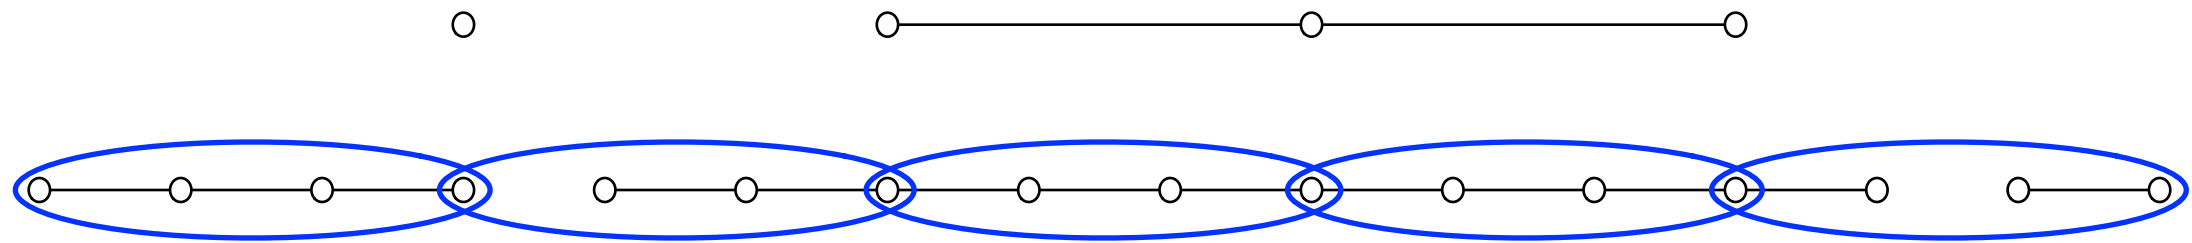
---



- Divide path into  $n/w$  subpath of  $w$  nodes.
- Level 1: The  $O(n/w)$  boundary nodes of the subpaths. Edge between two boundary nodes iff connected by subpath. Maintain using balanced relabeling data structure.
- Level 2: Short path data structure for each subpath.

# A Two-Level Solution

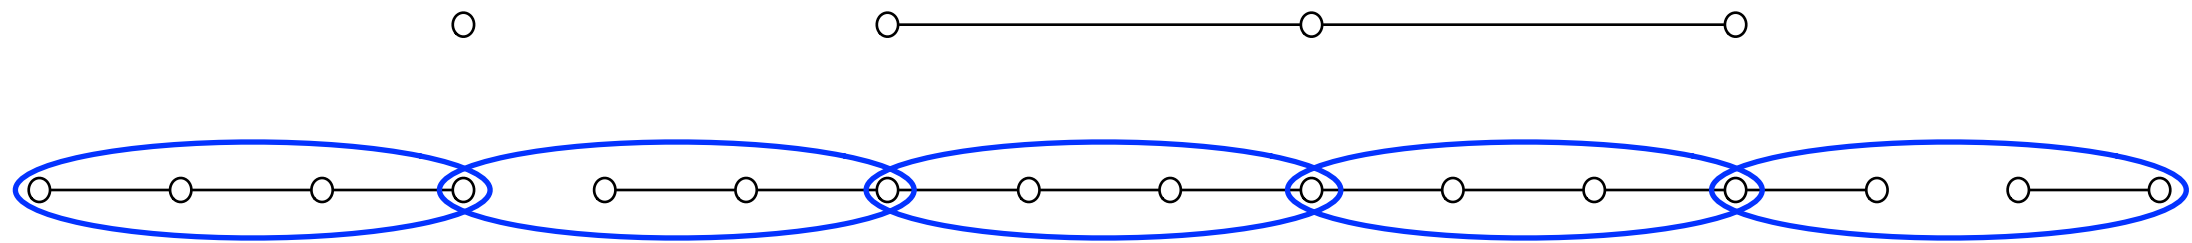
---



- `delete(e)`: delete edge in level 2. If first deletion in subpath also delete in level 1 using balanced relabeling strategy.
- `connected(v,u)`:
  - Case 1:  $v$  and  $u$  in same subpath. Use level 2 data structure.
  - Case 2:  $v$  and  $u$  in different subpaths. Return true iff
    - `connected(v, right-boundary(v))` &
    - `connected(u, left-boundary(u))` &
    - `connected(right-boundary(v), left-boundary(u))`

# A Two-Level Solution

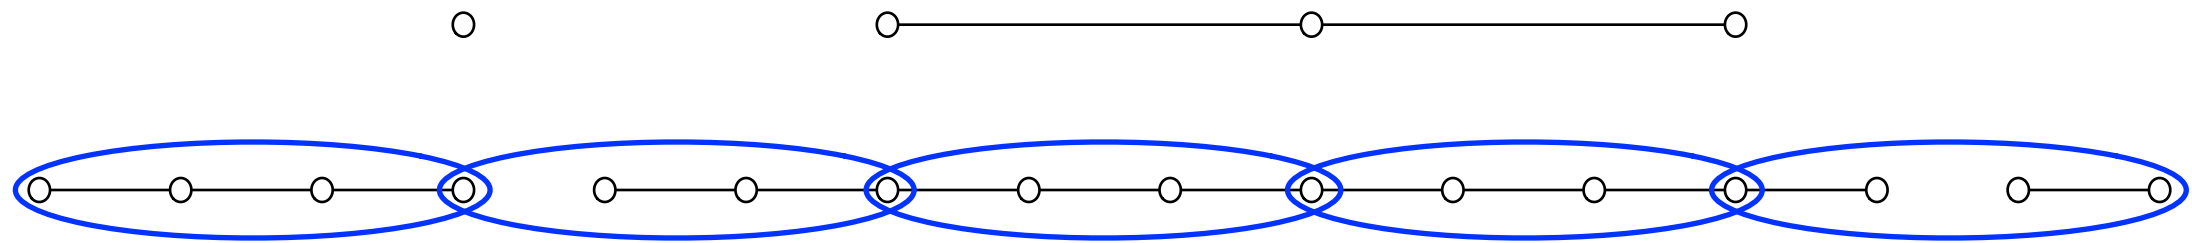
---



- $O(1)$  time for connected (at most 1 connected query in level 1 + 2 connected queries in level 2)
- $n-1$  delete operations:
  - Level 2:  $O(n)$
  - Level 1:  $O(n/w \cdot \log n) = O(n)$
- $\Rightarrow O(n)$  in total.

# Summary

---



- Theorem: We can solve decremental connectivity in paths in
  - $O(1)$  time for connected
  - $O(n)$  time for  $n-1$  deletions

# Decremental Connectivity in Trees



# Decremental Connectivity in Trees

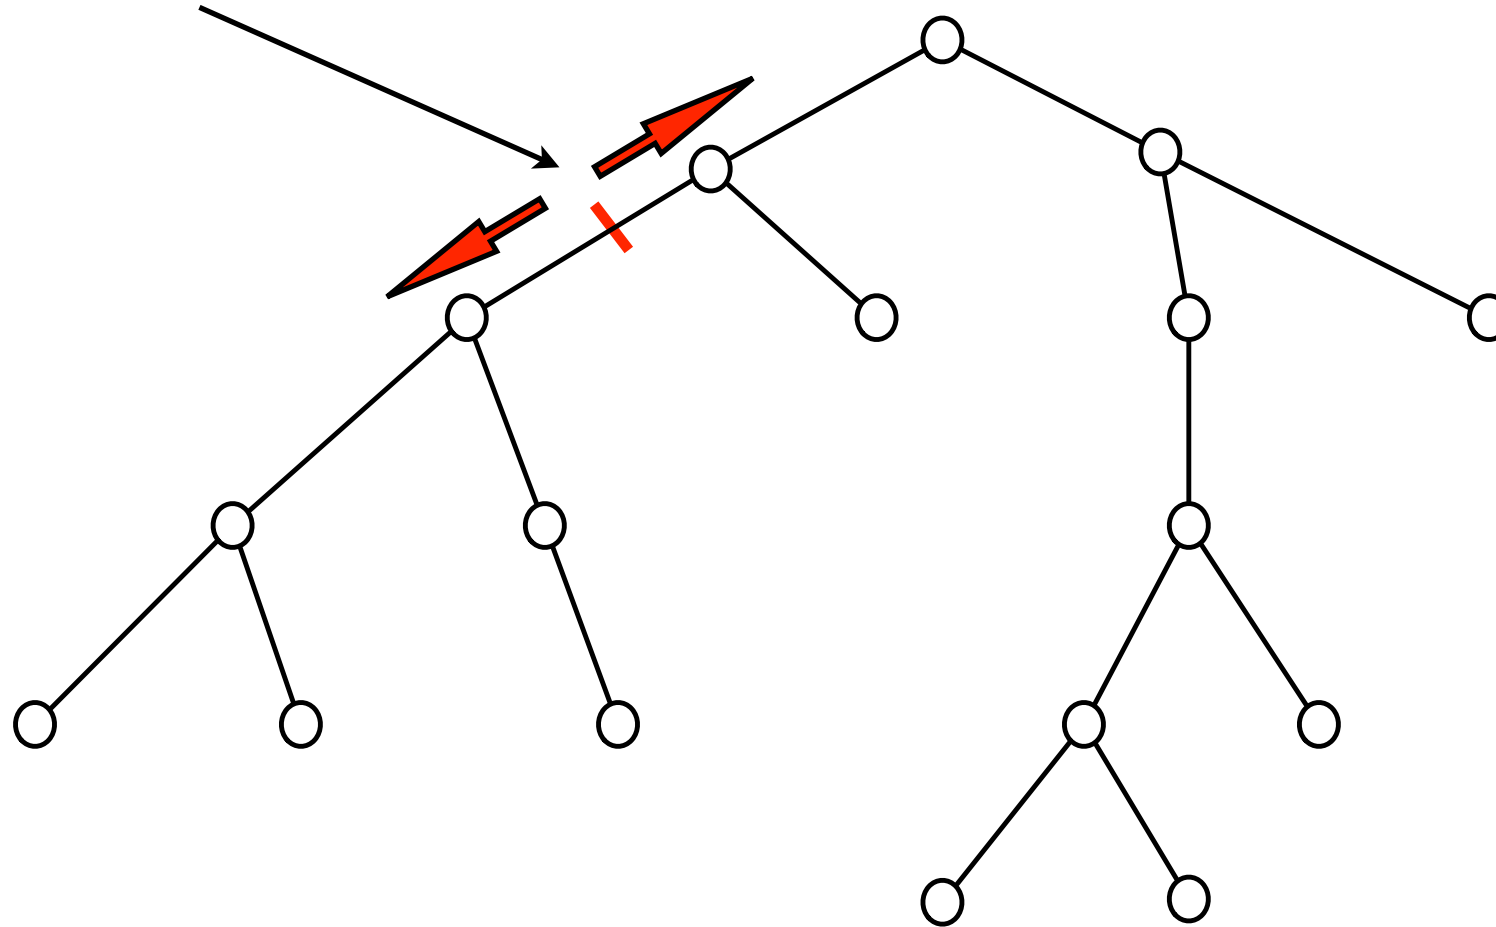
---

- Goal: Generalize two-level data structure from paths to trees.
- Simplifying assumption: Maximum degree of nodes in tree is 3.
- We need:
  - A balanced relabeling algorithm for trees.
  - An algorithm to divide trees into  $O(n/w)$  subtrees of  $\leq w$  nodes that only overlap in boundary nodes (does such division even exist?)
  - A fast data structure for decremental connectivity on subtrees with  $\leq w$  nodes.

# Balanced Relabeling for Trees

# Balanced Relabeling for Trees

Replace left and right search with Breadth-first search away from edge



# Balanced Relabeling for Trees

---

- Breadth-first search uses time linear in size of component.
- $\Rightarrow$  Same analysis as with paths
- $\Rightarrow O(n \log n)$  time for  $n-1$  deletes.

# Summary

---

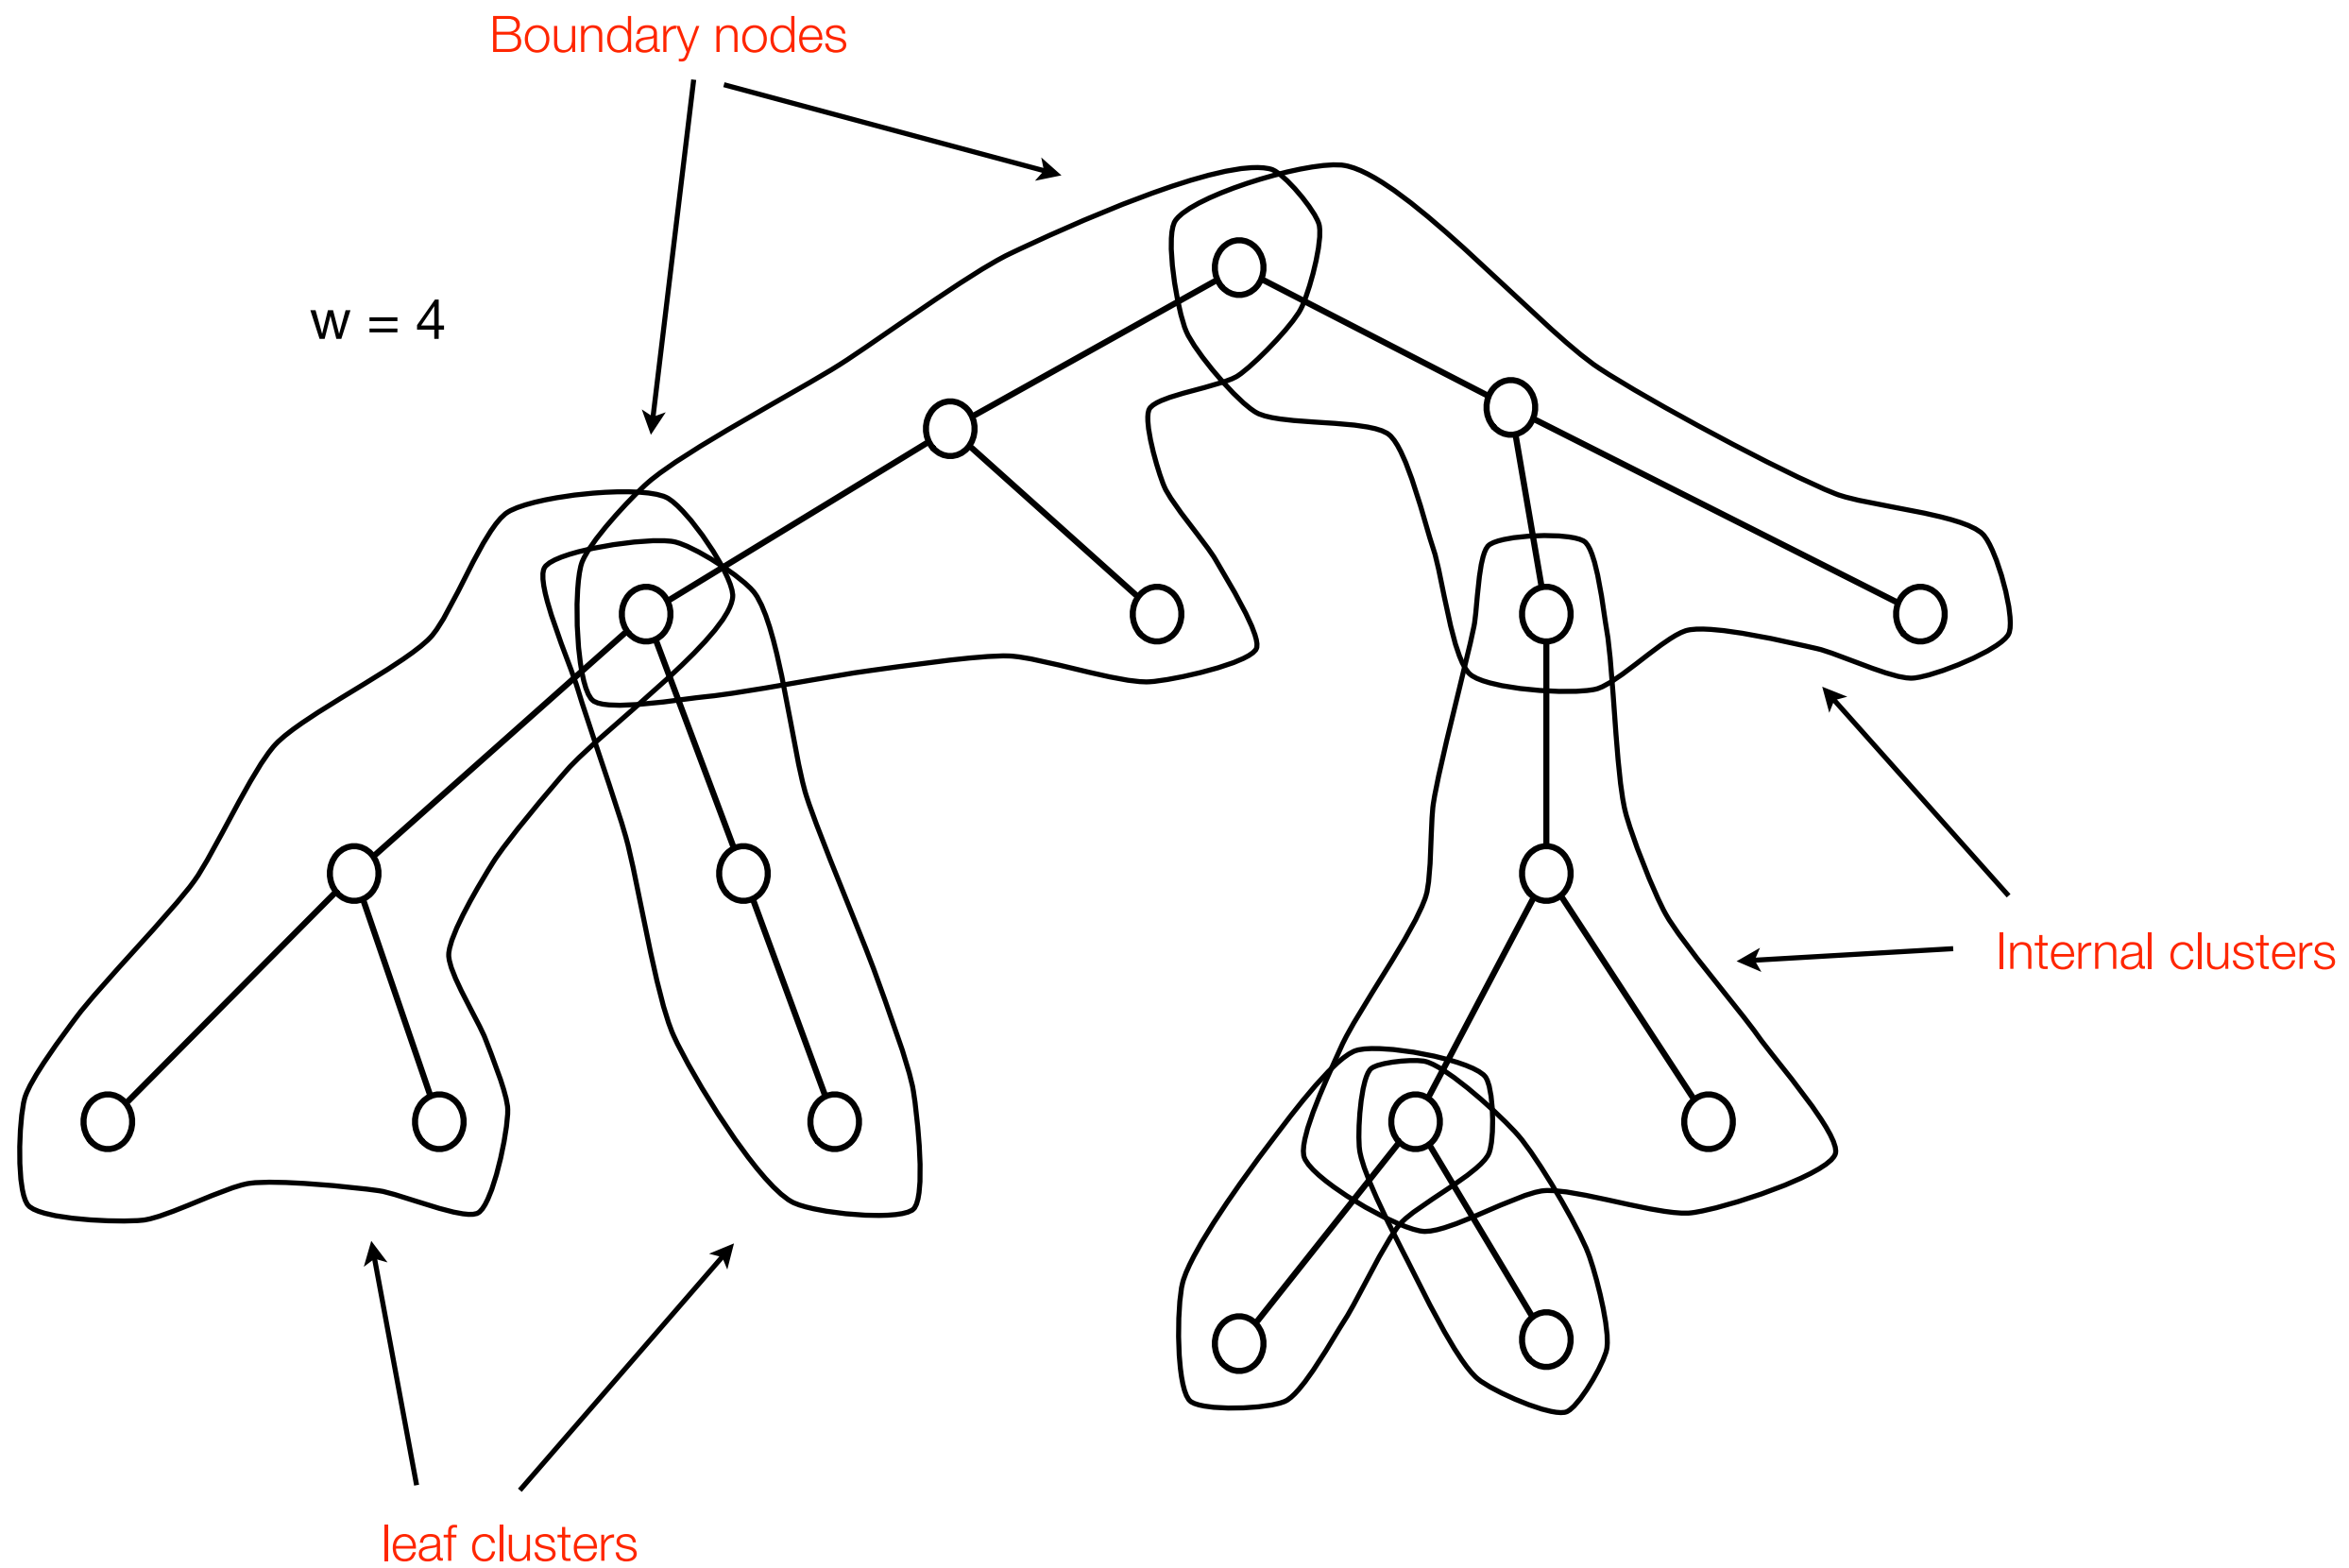
- Theorem: We can solve decremental connectivity in *trees* in
  - $O(1)$  time for connected.
  - $O(n \log n)$  time for  $n-1$  deletes.

# Tree Clustering

# Tree Clustering

---

- Goal: Given a tree  $T$  with maximum degree  $\leq 3$  compute a *cluster decomposition* of  $T$ :
  - Divide  $T$  into  $O(n/w)$  connected subtrees (*clusters*) of  $\leq w$  nodes.
  - Each cluster overlaps with other clusters in at most 2 *boundary nodes*.

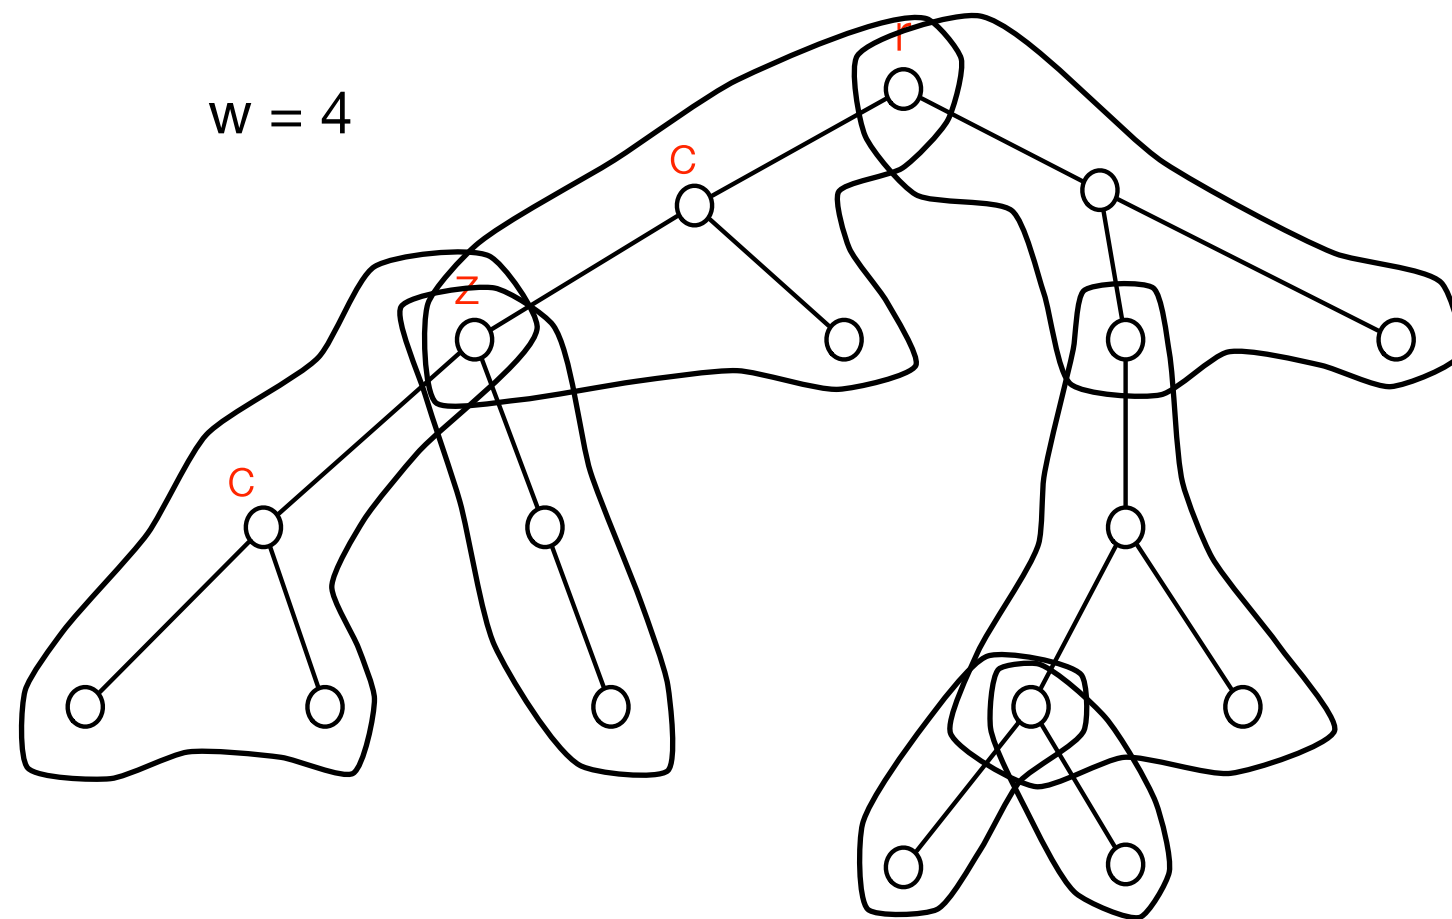




# Tree Clustering

---

- Lemma: A cluster decomposition exists and we can compute it in  $O(n)$  time.
- Main idea:
  - Root  $T$  at arbitrary node  $r$ .
  - Construct clusters greedily top-down.



- For each child  $c$  of root  $r$ :
  - If  $c$  has  $\leq w - 1$  descendants. Form leaf cluster from  $r$  and descendants of  $c$  with  $r$  as boundary node.
  - If  $c$  has  $\geq w$  descendants. Pick descendant  $z$  of max depth to form internal cluster of  $\leq w$  nodes with  $r$  and  $z$  as boundary nodes. Recurse on  $z$ .

# Tree Clustering

---

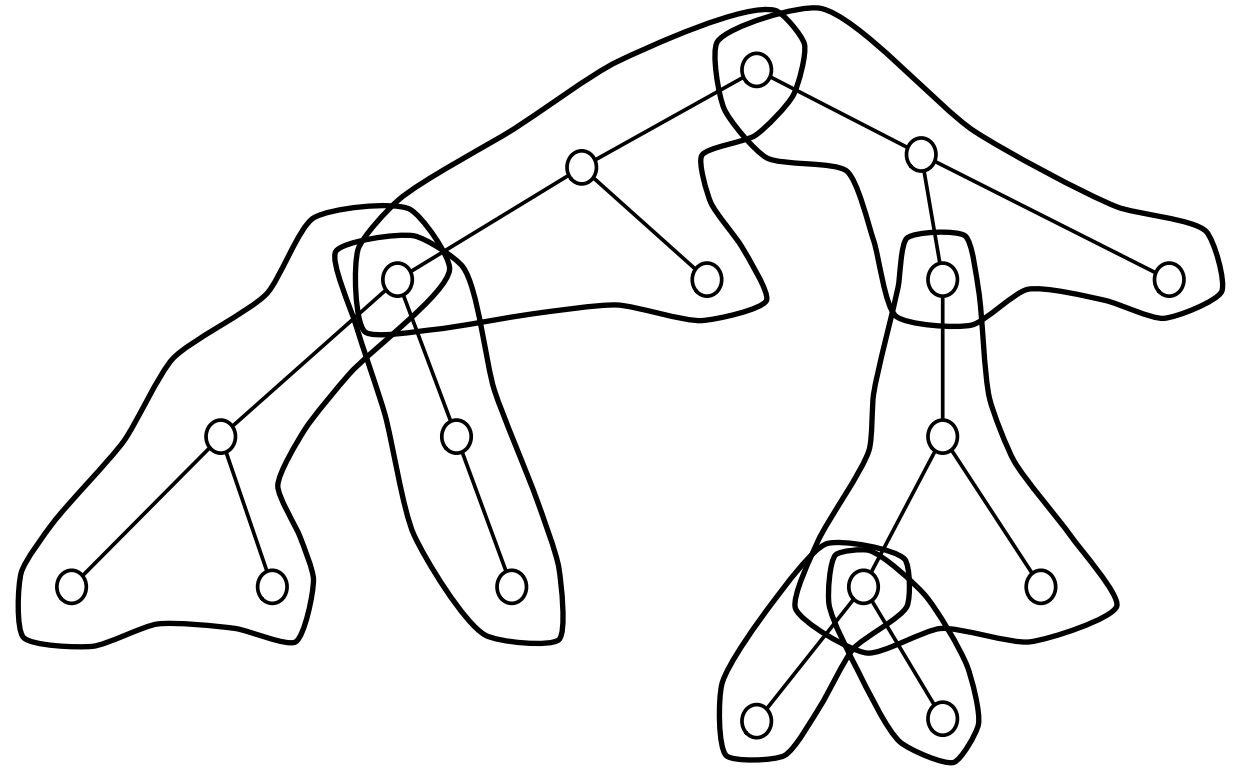
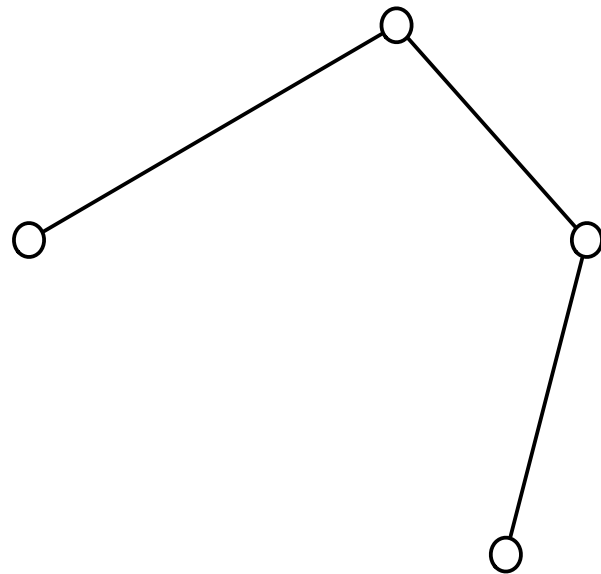
- Why does the clustering algorithm produce  $O(n/w)$  clusters each with  $\leq w$  nodes?
- Each cluster has  $\leq w$  nodes.
- How many clusters do we get?
- Intuition:
  - Greed  $\Rightarrow$  A constant fraction of cluster will have  $\Omega(w)$  nodes.
  - $\Rightarrow$  There will be at most  $O(n/w)$  clusters.

# Data Structure for Small Trees

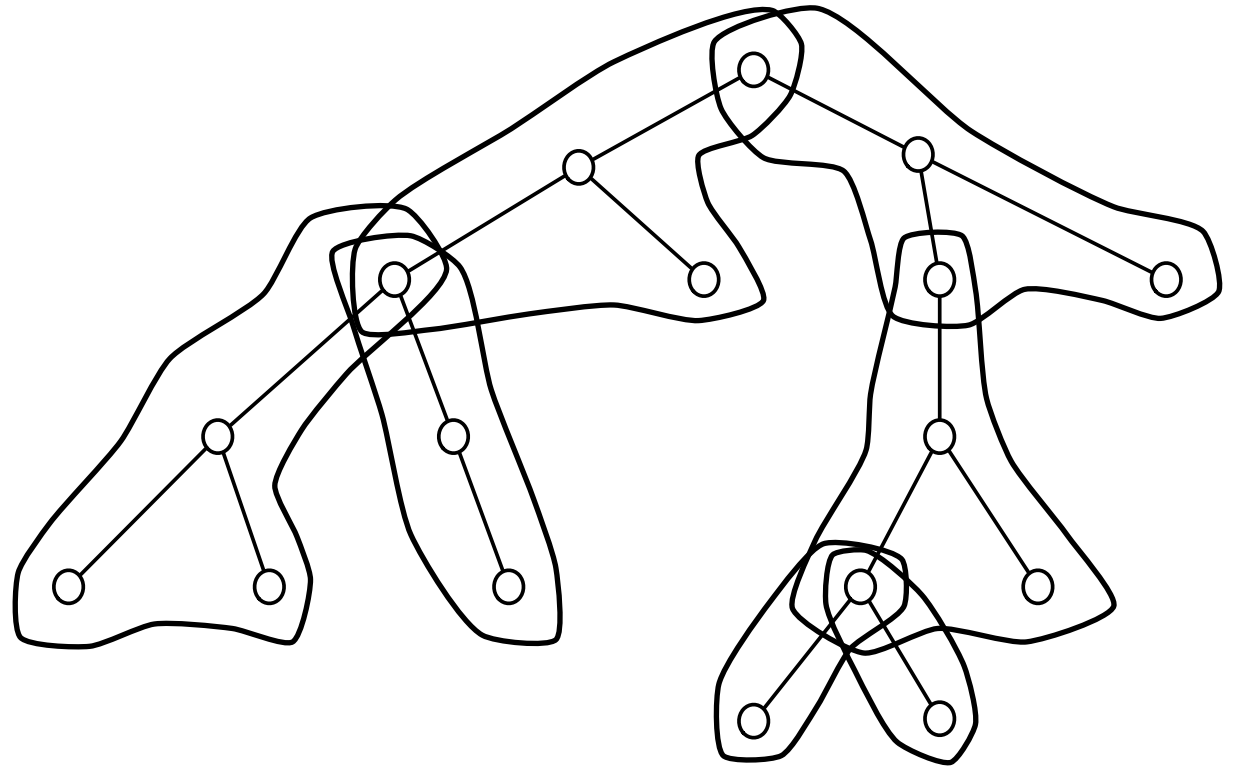
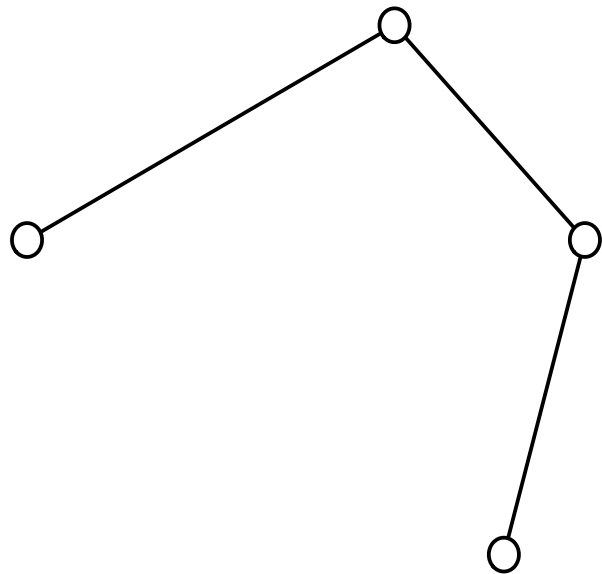
# Data Structure for Small Trees

---

- Theorem: We can solve decremental connectivity in trees with at most  $\leq w$  nodes in
  - $O(1)$  time for connected.
  - $O(1)$  time for delete



- Two-level data structure.
- Level 1: The  $O(n/w)$  boundary nodes of the clusters. Edge between two boundary nodes iff connected in cluster. Maintain using balanced relabeling data structure.
- Level 2: Small tree data structure for each cluster.



- Analysis: Exactly as in the path case
- Theorem: We can solve decremental connectivity in trees in
  - $O(1)$  time for connected
  - $O(n)$  time for  $n-1$  deletions

# Summary

---

- Decremental Connectivity in Trees Problem
- Decremental Connectivity in Paths
  - First Tradeoffs
  - Two-Level Solution
- Decremental Connectivity in Trees
  - Two-Level Solution



# References

---

- S. Alstrup, J. P. Secher, M. Spork: Optimal On-Line Decremental Connectivity in Trees, Inf. Process. Lett., 1997
- S. Alstrup, J. P. Secher, M. Thorup: Word encoding tree connectivity works. SODA, 2000